Quality for fitnessfirst.de

Find the expensive mistakes first, and find them early.

My starting point

HOW I PREPARED

- Explored the website as a visitor to understand the flows
- Checked if it feels slow as a user, or unusually slow (PageSpeed Insights, mobile)
- Looked at whether the site is responsive and ready for mobile
- Looked at accessibility: only a quick automatic scan (Google's tool), no manual check
- Nothing submitted, no login, no tests on your live site

WHAT I ASSUMED

- You have test environments next to production
- Payments are tested only in the provider's test mode, never on the live site (real money)
- Most users use a phone, not a laptop

How I explore

Exploratory testing: short sessions with a clear goal.

Example goal "Explore the checkout by changing my mind often, to find wrong totals."

WHAT I WRITE DOWN EVERY TIME

- Goal and setup (device, browser, which version of the page in an A/B test)
- What I covered
- Bugs, with steps and screenshots
- Open questions
- What I did NOT test

Saved next to the ticket, so the team can see it.

What I learned: the areas that matter most

Risk = how bad if it breaks x how likely it breaks

1. Membership checkout
2. Cookie consent and tracking
3. Pause, cancel, withdrawal
4. Trial training booking
5. Campaigns and offers
6. Club search
7. Member login
8. Course plan
9. Magazine, career

My plan for each risk

Area	Main risk	My plan
Membership checkout	Wrong price, or a contract that never arrives	Price rule tests, end-to-end tests on staging, automatic run on every push and every night
Cookie consent and tracking	Marketing trackers load after "decline"	Automatic check after "decline all", review with the data protection officer
Pause, cancel, withdrawal	A legal flow stops working	End-to-end test on every release, texts checked by hand
Trial training booking	Slot does not exist, booking never reaches the club	One end-to-end booking per phone type, check that it arrives
Campaigns and offers	An old offer still shows after it ended	Date tests, look at what the cache shows

The other areas lighter checks, mostly by hand.

Flow 1: Membership checkout, the risks

5 STEPS

Club → Plan → Extras → Your details → Payment

WHAT CAN GO WRONG

- Wrong price: club, plan, campaign, age, country
- Summary does not match the choice
- Customer gets stuck and leaves
- Contract or payment does not reach your systems
- Legal text missing or wrong

Most expensive a wrong price or a contract that was never created.

Flow 2: Trial training booking, the risks

3 STEPS

Club → Date and time → Your details

WHAT CAN GO WRONG

- Booked slot does not exist at the club
- Booking never reaches the club
- Date mistakes: today, weekend, month change, holidays
- Slow to react when someone taps (real users, mobile)

Automate slot rules, booking arrives at the club, one end-to-end booking per phone type.

Devices, users, outside systems

Devices: desktop browser first, then the mobile browser view. Real phones: iPhone Safari, Android Chrome.

Screen widths 360, 390, 768, 1280 pixels.

Users new visitors, students, people switching gyms, visitors from Austria, keyboard and screen reader users, slow mobile internet, people who decline cookies.

OUTSIDE SYSTEMS I SAW

- Usercentrics (cookie consent)
- VWO (A/B tests: not everyone sees the same page)
- Google Tag Manager and about 15 ad services
- Probably Baqend and Cloudflare (speed and caching)
- Unknown to me: payment, booking, member system

Checkout tests: what each one expects

Test pyramid: many small fast tests at the bottom, few slow ones at the top.

Exploratory (manual): new offers, odd combinations (student, other club), Austria. Expected: no wrong price or text.

End-to-end (few): pick club and plan, reach payment, on iPhone and Android. Change club mid-way. Expected: every step opens, price updates.

Visual and accessibility: Expected: looks right, keyboard reaches every step, errors are read out.

API and integration: plans, prices, contract, payment test mode. Expected: right data, contract arrives.

Unit (many): every price rule and campaign date. Expected: same input, same exact price.

What I automate, and what I don't

AUTOMATE WHEN

- it runs often, matters a lot, and the expected result is clear
- it is stable and cheap to maintain
- Example: price rules (unit tests)

KEEP MANUAL WHEN

- it is new or still changing
- it needs human judgement: look, wording, screen reader
- Examples: a new campaign page, the screen reader on checkout

WHAT I WOULD NOT TEST

- Outside services themselves: payment provider, cookie tool, ad services. I test only our side.
- Load and security attacks on the live site
- Low-risk areas in depth, like the magazine: a quick look is enough

The quality process I would encourage

Six phases, one question: what could go wrong, and how do we prevent it?

Bugs cost less when we find them early.

1
Requirements
understand the why

2
Risk
decide together what matters most

3
Development
explore and pair early

4
Automation
automate the right things

5
Release
done is a criterion, not a feeling

6
After release
learn from what got through

Quality is the team's job

Product	clear goals and acceptance criteria. Decides on risk together with the team.
Developers	unit and API tests for their code. Fix what breaks the pipeline.
Business departments	check that it is right for the business (acceptance).
QA	gate keeper at release, and a resource for everyone: • helps write acceptance criteria • pairs with developers • keeps the test cases alive • coaches the team on testing

From idea to live website



Feature	Push and merge request	Dev	Stage	Release
acceptance criteria and risks agreed	unit and API tests run on every push. They must pass before the code reaches master.	integration tests. Smoke test after each deployment.	smoke test after each deployment, then end-to-end regression, visual, accessibility, page speed, consent check, acceptance by business. Plus a full run every night.	all green, then live, with a smoke test after the deployment. Rollback ready.

Test data no real customer data on Stage. Copied data is anonymised.

When do we stop a release?

STOPS THE RELEASE

- Wrong price, checkout or booking broken
- Cancel or pause broken
- Cookie decline does not work
- Security problem, or someone sees another person's data
- Main flow not usable with a keyboard

CAN GO, WITH CONDITIONS

- Small bug with an easy workaround, ticket and fix date
- Problem on a low-traffic page
- New feature switched off with a feature flag

Unreliable tests run once more. If it passes, it goes to quarantine with an owner and a deadline. Never just ignored.

Three more areas: accessibility, privacy, speed

Area	Automated	Manual	Limit (gate)
Accessibility	Accessibility scan on every merge request	Keyboard and screen reader on main flows every release, expert audit once a year	No critical issues on checkout and booking
Data protection and consent	After "decline all", only allowed services load	Check banner texts with the data protection officer	Any marketing tracker after decline stops the release
Page speed	Lighthouse in the pipeline, timing of booking steps	Look at real-user data every week	Real users: content in 2.5 s, reaction to a tap in 200 ms. Pipeline: no page slower than today

How can AI help us, and where?

AI HELPS WITH

1. **Requirements:** finds gaps and missing edge cases in acceptance criteria
2. **Exploratory testing:** agents explore with a goal and must show evidence for every result
3. **Failed tests:** groups failures, suggests "real bug" or "flaky"

ALWAYS WITH PEOPLE

- Final priority of bugs
- Release or stop
- Legal, accessibility and data protection decisions

Rules no real customer data in AI tools. Every result needs evidence. AI must list what it did not test.



Extra slides

Page speed details

Page (mobile)	Real users: LCP	Real users: INP	Real users: CLS
Homepage	2 s	208 ms	0.03
Membership	1.1 s	227 ms	0.04
Trial training	1.3 s	361 ms	0.05

Limits for good (75% of visits): LCP 2.5 s, INP 200 ms, CLS 0.1. Source: PageSpeed Insights, 27 October 2026.

Test data and environments

- Stage: anonymised copy of production, or generated test data
- Fixed test accounts and test clubs, reset regularly
- Payment provider in test mode, never real payments
- Outside systems we cannot control: simulated (mocked) in API tests, real only in a few end-to-end tests on Stage
- Test emails only on reserved test domains

My AI testing agents

- 4 agents: explorer, component, flow, API
- Rules: only allowed test sites, no real purchases, secrets hidden, evidence for every result
- Results: pass, fail with priority, or "suspicious" with a reason
- Several runs are merged, because AI results vary from run to run
- A top-priority failure (P1), like a broken checkout, can stop the pipeline. Humans decide on the rest.

My first weeks

Weeks 1 to 2	learn the product, the team and the release process. Read the bug history.
Weeks 3 to 4	risk map with the team, first quick wins (consent check, smoke tests)
Month 2 to 3	regression suite for checkout and booking, quality gates in the pipeline

One normal week

Planning	Three Amigos on new stories, risks agreed
Every day	pairing, merge request checks, quick questions
During the sprint	exploratory sessions on new features
Before release	regression on Stage, acceptance with business
Release day	check that all gates are green, then release
After release	look at what got through, add a test for it